

Large-scale Distributed Training for LLMs

Ayush Maheshwari

Sr. Solutions Architect, NVIDIA



Sessions

- | | |
|---|------------------|
| 1. Understanding the hardware | (30 mins) |
| a) GPU vs CPU | |
| b) GPU communication primitives | |
| c) System Topology | |
| 2. Large scale data curation for LLM training | (1 hour) |
| a) Deep-dive into aspects of data curation | |
| b) Mixed-precision training | |
|
BREAK

 | |
| 3. Distributed and stable LLM training on a large-scale cluster | (10 mins) |
| a) Parallelism techniques | |
| b) Frameworks and wrappers | |
| c) Recipes and best practices | |
| 4. Inference | (15 mins) |
| a) Inference with build.nvidia.com | |
| b) Synthetic data generation | |


Register for GTC 2026

<https://tinyurl.com/nvgtc2026>



Scan QR

Language Model became more complex and larger



Transformer Architecture




Figure 1: The Transformer - model architecture.

Transformer Architecture






Figure 1: The Transformer - model architecture.

Transformer Architecture







Figure 1: The Transformer - model architecture.



A large, abstract graphic on the left side of the slide features several curved, overlapping planes in shades of lime green, yellow, and dark green. The planes are arranged in a way that suggests depth and perspective, creating a sense of a complex, multi-layered structure.

Agenda


1. Concepts of Parallelism – Data, Tensor and Pipeline
2. How does these work together ?
3. Mixed –precision training

Not Covering

Transformer architecture, FP8 training

Training a Neural Network


Single GPU



1. Read the data
2. Transport the data
3. Pre-process the data
4. Queue the data
5. Transport the data
6. Calculate activations for layer one
7. Calculate activations for layer two
8. Calculate the output
9. Calculate the loss
10. Backpropagate through layer three
11. Backpropagate through layer two
12. Backpropagate through layer one
13. Execute optimization step
14. Update the weights
15. Return control

Training a Neural Network


Multiple GPUs with DDP




Distributed Training: Parallelism

Parallelism : An overview

Data parallelism (DP)




Model parallelism (MP)



n copies of model parameters

Single copy of model parameters

Data parallelism



- Naïvely, model copy on each GPU
- Reductions of weight gradients at the end of every iteration to coalesce updates across replicas
- Our data parallelism implementation involves a simple DDP wrapper, with largely the same interface as PyTorch's DDP

Data Parallelism (Distributed Optimizer)

High Level Abstraction

Do {

Forward Path (activations) – calculate error


Backward Path – calculate gradients

Reduce Scatter gradients – each source gets a different part of the results

Update network weights (a.k.a. optimizer) **available local part of the gradient weights**

All Gather weights calculated in each GPU

} While error is above threshold / not decreasing anymore



Data Parallelism (Distributed Optimizer)

High Level Abstraction

Do {

Forward Path (activations) – calculate error


Backward Path – calculate gradients

Reduce Scatter gradients – each source gets a different part of the results

Update network weights (a.k.a. optimizer) **available local part of the gradient weights**


All Gather weights calculated in each GPU

} While error is above threshold / not decreasing anymore




Distributed Data Parallel - DDP

PyTorch: Streamline API for Fully Sharded Data Parallel (FSDP)




per-GPU throughput vs number of GPUs



Sharded Data Parallelism


ZeRO: Zero Redundancy Optimizer

- ZeRO removes the redundancy across data parallel process
- Partitioning optimizer states, gradients and parameters (3 stages) for a progressive memory savings and Communication Volume




Sharded Data Parallelism

ZeRO: Stage 1




Sharded Data Parallelism

ZeRO: Stage 2



Sharded Data Parallelism

ZeRO: Stage 3



GPU Memory occupation

Let's review what is in your GPU memory

- **Model Weights**

- 4 bytes * number of parameters for fp32 training
- 6 bytes * number of parameters for mixed precision training (maintains a model in fp32 and one in fp16 in memory)

- **Optimizer States**

- 8 bytes * number of parameters for normal AdamW (maintains 2 states)
- 2 bytes * number of parameters for 8-bit AdamW optimizers like bitsandbytes
- 4 bytes * number of parameters for optimizers like SGD with momentum (maintains only 1 state)

- **Gradients**

- 4 bytes * number of parameters for either fp32 or mixed precision training (gradients are always kept in fp32)

- **Forward Activations**

- size depends on many factors, the key ones being sequence length, hidden size and batch size

Distributed optimizer to reduce memory

$$\text{Number of bytes of state per parameter} = 2 + 4 + 4 + 4 + 4$$




bf16 params
fp32 grads

fp32 copy of params
fp32 Adam states

$$\text{Number of bytes of state for Nemotron-4 340B model} = 18 \cdot 340\text{B} = \mathbf{6120 \text{ GB}}$$

Redundant optimizer state over DP replicas can be partitioned
fp32 gradient all-reduces →
fp32 gradient reduce-scatters + **bf16** param all-gathers

Data Parallelism / Model Parallelism



Data is too large, accelerated by processing in parallel


Model is too large, cannot fit in a single device / machine

TECHNOLOGIES THAT ENABLE SCALING LARGE MODELS

Complementary Types of Parallelism


Pipeline (Inter-Layer) Parallelism

Split contiguous sets of layers across multiple GPUs
Layers 0,1,2 and layers 3,4,5 are on different GPUs
Maximizes GPU utilization in single-node



Tensor (Intra-Layer) Parallelism


Split individual layers across multiple GPUs
Both devices compute different parts of Layers 0,1,2,3,4,5
Minimizes Latency in single-node



TENSOR PARALLELISM

Why?

- Relatively simple to implement
- Easier to load-balance
- Less restrictive on the batch-size (avoids bubble issue in pipelining)
 - Tensor parallelism is orthogonal to pipeline parallelism: very large models such as GPT-3 use both
- NVIDIA DGX servers with NVSwitch
 - DGX A100 has 600 GB/sec GPU-to-GPU bidirectional bandwidth
- Tensor parallelism works well for large matrices
 - Example: Transformers have large GEMMs



MLP TENSOR PARTITIONING

Focus on the GeLU operation:

- Approach 1: Split X column-wise and A row-wise:

$$X = [X_1, X_2] \quad A = \begin{bmatrix} A_1 \\ A_2 \end{bmatrix} \quad \rightarrow \quad Y = \text{GeLU}(X_1A_1 + X_2A_2)$$


- Before GeLU we will need a communication point
- Approach 2: Split A column-wise:

$$A = [A_1, A_2] \quad \rightarrow \quad [Y_1, Y_2] = [\text{GeLU}(XA_1), \text{GeLU}(XA_2)]$$

- No communication is required

$$Z = \text{Dropout}(YB)$$

$$Y = \text{GeLU}(XA)$$



MLP TENSOR PARTITIONING

Focus on the GeLU operation:

- Approach 1: Split X column-wise and A row-wise:

$$X = [X_1, X_2] \quad A = \begin{bmatrix} A_1 \\ A_2 \end{bmatrix} \quad \rightarrow \quad Y = \text{GeLU}(X_1A_1 + X_2A_2)$$


- Before GeLU we will need a communication point


- Approach 2: Split A column-wise:

$$A = [A_1, A_2] \quad \rightarrow \quad [Y_1, Y_2] = [\text{GeLU}(XA_1), \text{GeLU}(XA_2)]$$

- No communication is required


Chosen approach






MLP TENSOR PARTITIONING


GeLU Column Parallel






MLP TENSOR PARTITIONING

Dropout Row Parallel




$g :$
forward: $Z = Z_1 + Z_2$
backward: $\frac{\partial L}{\partial Z_i} = \frac{\partial L}{\partial Z}$ (identity)

$$(Y_1, Y_2) * \begin{pmatrix} B_1 \\ B_2 \end{pmatrix} \iff YB$$



TENSOR PARTITIONING

Each layer of model is partitioned over multiple GPUs





$g \rightarrow$ All-reduction ($Y_1B_1 + Y_2B_2$) in forward pass

Slow across inter-server communication links

f and g are conjugate, f is identity operator in the forward pass and all-reduce in the backward pass while g is all-reduce in forward and identity in backward.


SELF-ATTENTION TENSOR PARTITIONING

Self-Attention is more complex than MLP




TENSOR PARALLEL TRANSFORMER LAYER

All Together




PIPELINE PARALLELISM

Challenges




PIPELINE PARALLELISM

Challenges - Idle Workers




PIPELINE PARALLELISM

Split batch into micro batches and pipeline execution




Split batch into micro batches and pipeline execution




PIPELINE PARALLELISM

Split batch into micro batches and pipeline execution




PIPELINE PARALLELISM

Split batch into micro batches and pipeline execution




PIPELINE PARALLELISM

Split batch into micro batches and pipeline execution




PIPELINE PARALLELISM

Split batch into micro batches and pipeline execution



PIPELINE PARALLELISM


Split batch into micro batches and pipeline execution



Pipeline model parallelism

- Layers / operators in model sharded over GPUs (i.e., each GPU is responsible for a subset of layers in the model)
- Each batch split into smaller microbatches and execution pipelined across these microbatches
- Point-to-point communication between consecutive pipeline stages
- Pipeline bubble at the start and end of every batch (equal to $(p - 1)$ microbatches' forward and backward passes)

Interleaved pipeline parallelism



Smaller pipeline bubble but more communication

Mixed Precision Training

Benefits of less bits

Memory

Weights and tensors occupy less space in memory

Bandwidth

Faster data movement from main memory HBM to cores (and vice versa)



Compute

More TFLOPS with less bits
Faster matrix multiplications


	sign	exponent					mantissa										
FP16	0	0	1	1	0	1	1	0	0	1	0	1	0	0	1	1	= 0.395264
BF16	0	0	1	1	1	1	1	0	1	1	0	0	1	0	1	0	= 0.394531
FP8 E4M3	0	0	1	0	1	1	0	1									= 0.40625
FP8 E5M2	0	0	1	1	0	1	1	1	0								= 0.375

Wider distributions suffer more with Quantization

Gradients have the widest distributions




Scaling Factors to Keep Tensors within Range



Automatic Mixed Precision (AMP)

Concept

- Maintain a primary copy of weights in FP32.
- Initialize scaling factor S to a large value.
- For each iteration:
 - Make an FP16 copy of the weights.
 - Forward propagation (FP16 weights and activations).
 - Multiply the resulting loss with the scaling factor S .
 - Backward propagation (FP16 weights, activations, and their gradients).
 - If there is an Inf or NaN in weight gradients:
 - Reduce S .
 - Skip the weight update and move to the next iteration.
 - Multiply the weight gradient with $1/S$.
 - Complete the weight update (including gradient clipping, etc.).
 - If there hasn't been an Inf or NaN in the last N iterations, increase S .



Towards finer Scaling Factors

Ampere

FP16 mixed precision
one scale for all tensors


Hopper

FP8 training
one scale per tensor


Blackwell

MXFP8 training*
one scale per block of 32 elements

FP8



MXFP8



Scaling factor



Scaling factor

E8MO	E8MO

*MX stands for microscaled formats, read [OCP spec](#)


GEMM operation with MXFP8 operands

Operand A = { a_i, SF_A }


Operand B = { b_i, SF_B }

$$Output = (SF_A \times SF_B) \sum_0^{3^1} a_i \times b_1$$

Scales




Quantized values (FP8)




Operand A

Scales



Quantized values (FP8)




Operand B

FP32 Output

Tensors are quantized to MXFP8 after FP32 accumulation

How to cast back to MXFP8?



Transformer Engine

- An open-source library implementing the FP8 recipe for Transformer building blocks
- Optimized for FP8 and other datatypes
- PyTorch and JAX are supported frameworks
- Composable with the native framework operators
- Supports different types of model parallelism
 - DP, TP, PP, CP
- cuDNN kernels available such as GroupedGEMM
- <https://github.com/NVIDIA/TransformerEngine>
- Docs:
 - <https://docs.nvidia.com/deeplearning/transformer-engine/user-guide/index.html>

```
import torch
import transformer_engine.pytorch as te
from transformer_engine.common import recipe

# Set dimensions.
in_features = 768
out_features = 3072
hidden_size = 2048

# Initialize model and inputs.
model = te.Linear(in_features, out_features, bias=True)
inp = torch.randn(hidden_size, in_features, device="cuda")

# Create MXFP8 recipe.
fp8_recipe = recipe.MXFP8BlockScaling()

# Enable autocasting to FP8.
with te.fp8_autocast(enabled=True, fp8_recipe=fp8_recipe):
    out = model(inp)

# Calculate loss and gradients.
loss = out.sum()
loss.backward()
```

New FP8 Recipes for Blackwell and Hopper



TransformerEngine

[Hopper] Current Scaling + 1st & last layer BF16

- Current scaling is more stable than delayed scaling, but a bit slower
- Keeps the more sensitive 1st and last layer in BF16
- E4M3 for weights and activations, E5M2 for gradients

[Hopper] NV Subchannel Recipe (DeepSeek-V3 like)

- As DeepSeek-V3 pretraining
- 1x128 blocks for input and output_grad, 128x128 blocks for weights
- E4M3 for all weights, acts, and grads

[Blackwell] MXFP8 Blockwise Scaling

- Different scaling factor for each block of 32 values in a tensor
- E4M3 for all weights, acts, and grads

build.nvidia.com

API for accessing open models

Quick Start Select a model and click view code to make an API request in minutes

meta / llama-4-maverick-17b-128e-instruct

View Code

API Keys View & manage your API Keys

Your API Rate Limit Up to 40 rpm

Total 1 Inactive 0 Manage API Keys

Discover AI Agents Exclusively Available for NVIDIAians

View All →

Explore a wide range of AI Agents designed to enhance productivity and streamline workflows.




agent

NVINFO-AI

NVINFO AI is a conversational AI agent designed to help find answers from company knowledge.

internal-agent multi-agent rag +5




agent

BugNemo AI Search

AI Bug Assistant enhances code review efficiency through automated bug analysis.

internal-agent rag summarization +3



agent

CodeCritic

Generative AI based automated code review based on code changes made in SCMs(Gitlab, Gerrit and...

internal-agent code generator +2



agent

Supply Chain AI Planner

Real-time optimization and scenario global logistics operations

internal-agent multi-agent

Workshop Feedback

<https://tinyurl.com/nv-feedback>



Thank You