

Unlocking AI Performance with NeMo Curator: Scalable Data Processing for LLMs

Ayush Maheshwari
Sr. Solutions Architect, NVIDIA

<https://github.com/ayushbits/llm-development>

ayushbits.github.io

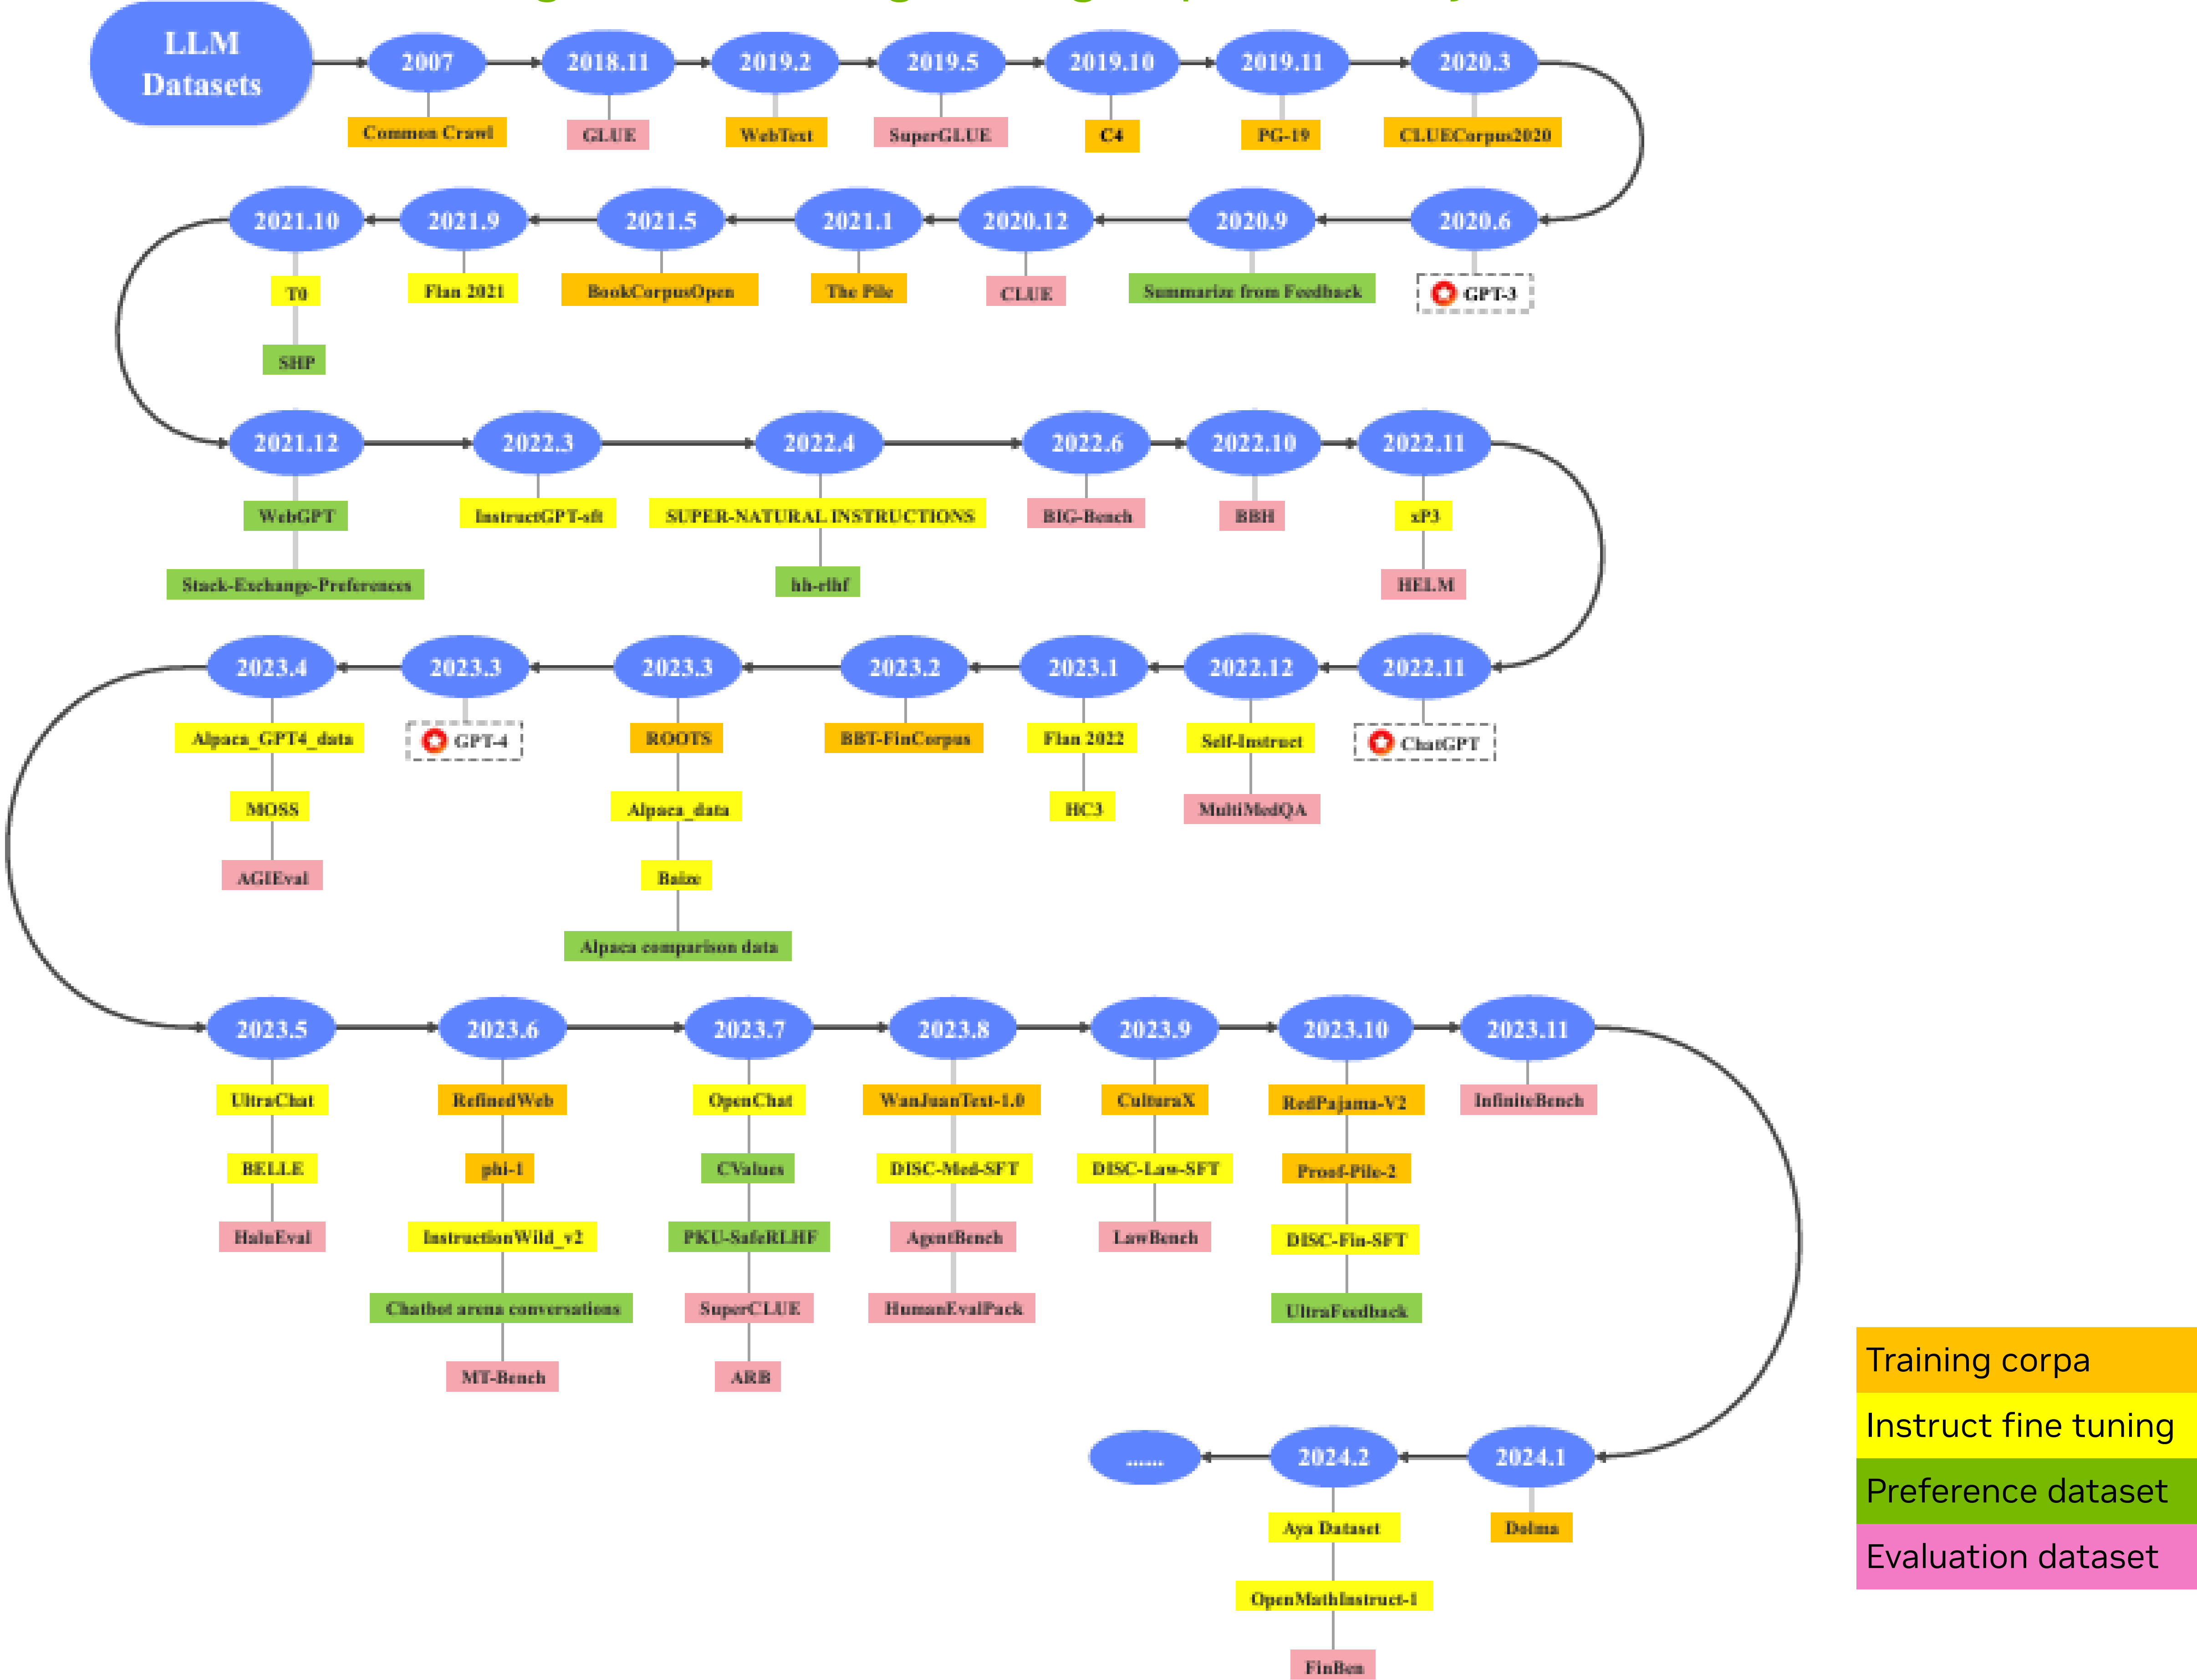


Sessions

1. Cluster health-check using NCCL, MLPerf, HPL **(1 hour) - Completed**
 - a) Understand the hardware and its performance on multiple GPUs.
 - b) Ensure that your training performance aligns with the h/w benchmarks
 - c) Evaluate the cluster to ensure platform fits within your needs.
2. Large scale data curation for LLM training **(1 hour) - Today**
 - a) Deep-dive into aspects of data curation
 - b) Mixed-precision training
3. Distributed and stable LLM training on a large-scale cluster **(1.5 hour)**
 - a) Parallelism techniques
 - b) Frameworks and wrappers
 - c) Recipes and best practices
4. Post-training and evaluation of pre-trained LLM **(1 hour)**
 - a) Sync between training data and expected performance
 - b) Algorithms and frameworks
5. Fine-tuning and deployment **(1 hour)**
 - a) Dynamic and static batching, state management, inference server
 - b) Best practices for optimizing model


LLMs Are Trained on Internet Scale Data

Data generated is growing exponentially




Data Processing for LLMs Needs Accelerated Computing


Petabytes of Data Generated Yearly




LLMs Trained on Internet Scale Data




Moore's Law Has Ended




Data Processing for Different LLM Needs



Training Foundation Model



Fine-Tuning Foundation Model




Retrieval Augmented Generation (RAG)


Data Size	TB and PB	GBs	GBs
Compute Scale	Supercomputer	Single-node	Single GPU
Frequency	One-time	Iterative	Iterative & Continuous

Challenges with Existing Solutions for Training Foundation Models


Inefficiencies lead to higher TCO and slower time to market



Longer Processing Time



Un-optimized Models



Un-optimized Pipelines



Knowledge & Expertise

NeMo Curator - Overview

Scalable, configurable pipelines to curate text, image and video datasets lead to more accurate applications

Higher Accuracy



Improve accuracy with less data and less training compute

Faster Processing



GPU acceleration with RAPIDS for **Dedupe (Exact, Fuzzy, Semantic)** and **Quality Classifier Models**

Scalability




Up to **100+ PB data** by scaling across multiple nodes

Classifier Models



State-of-the-Art quality classifier NIM microservice for safety, content, and diversity

Deploy anywhere



Python APIs in a customizable and modular OSS library, runnable across CSP and On-prem

[GitHub](#)

[NeMo framework container](#)

[PyPI](#)

Microservice
(coming soon)




Text Processing

High Quality Data Processing Maximizes Model Performance

Data Curation helps build SOTA Models

LLM Accuracy Improvement on Curated Data



Why is Data Curation Important?

State-of-the-art data curation is essential for developing state-of-the-art models across all modalities

Higher Accuracy



Properly curated data leads to more improved accuracy across tasks

TCO Savings



Decreases both training and inference costs significantly

Faster Training



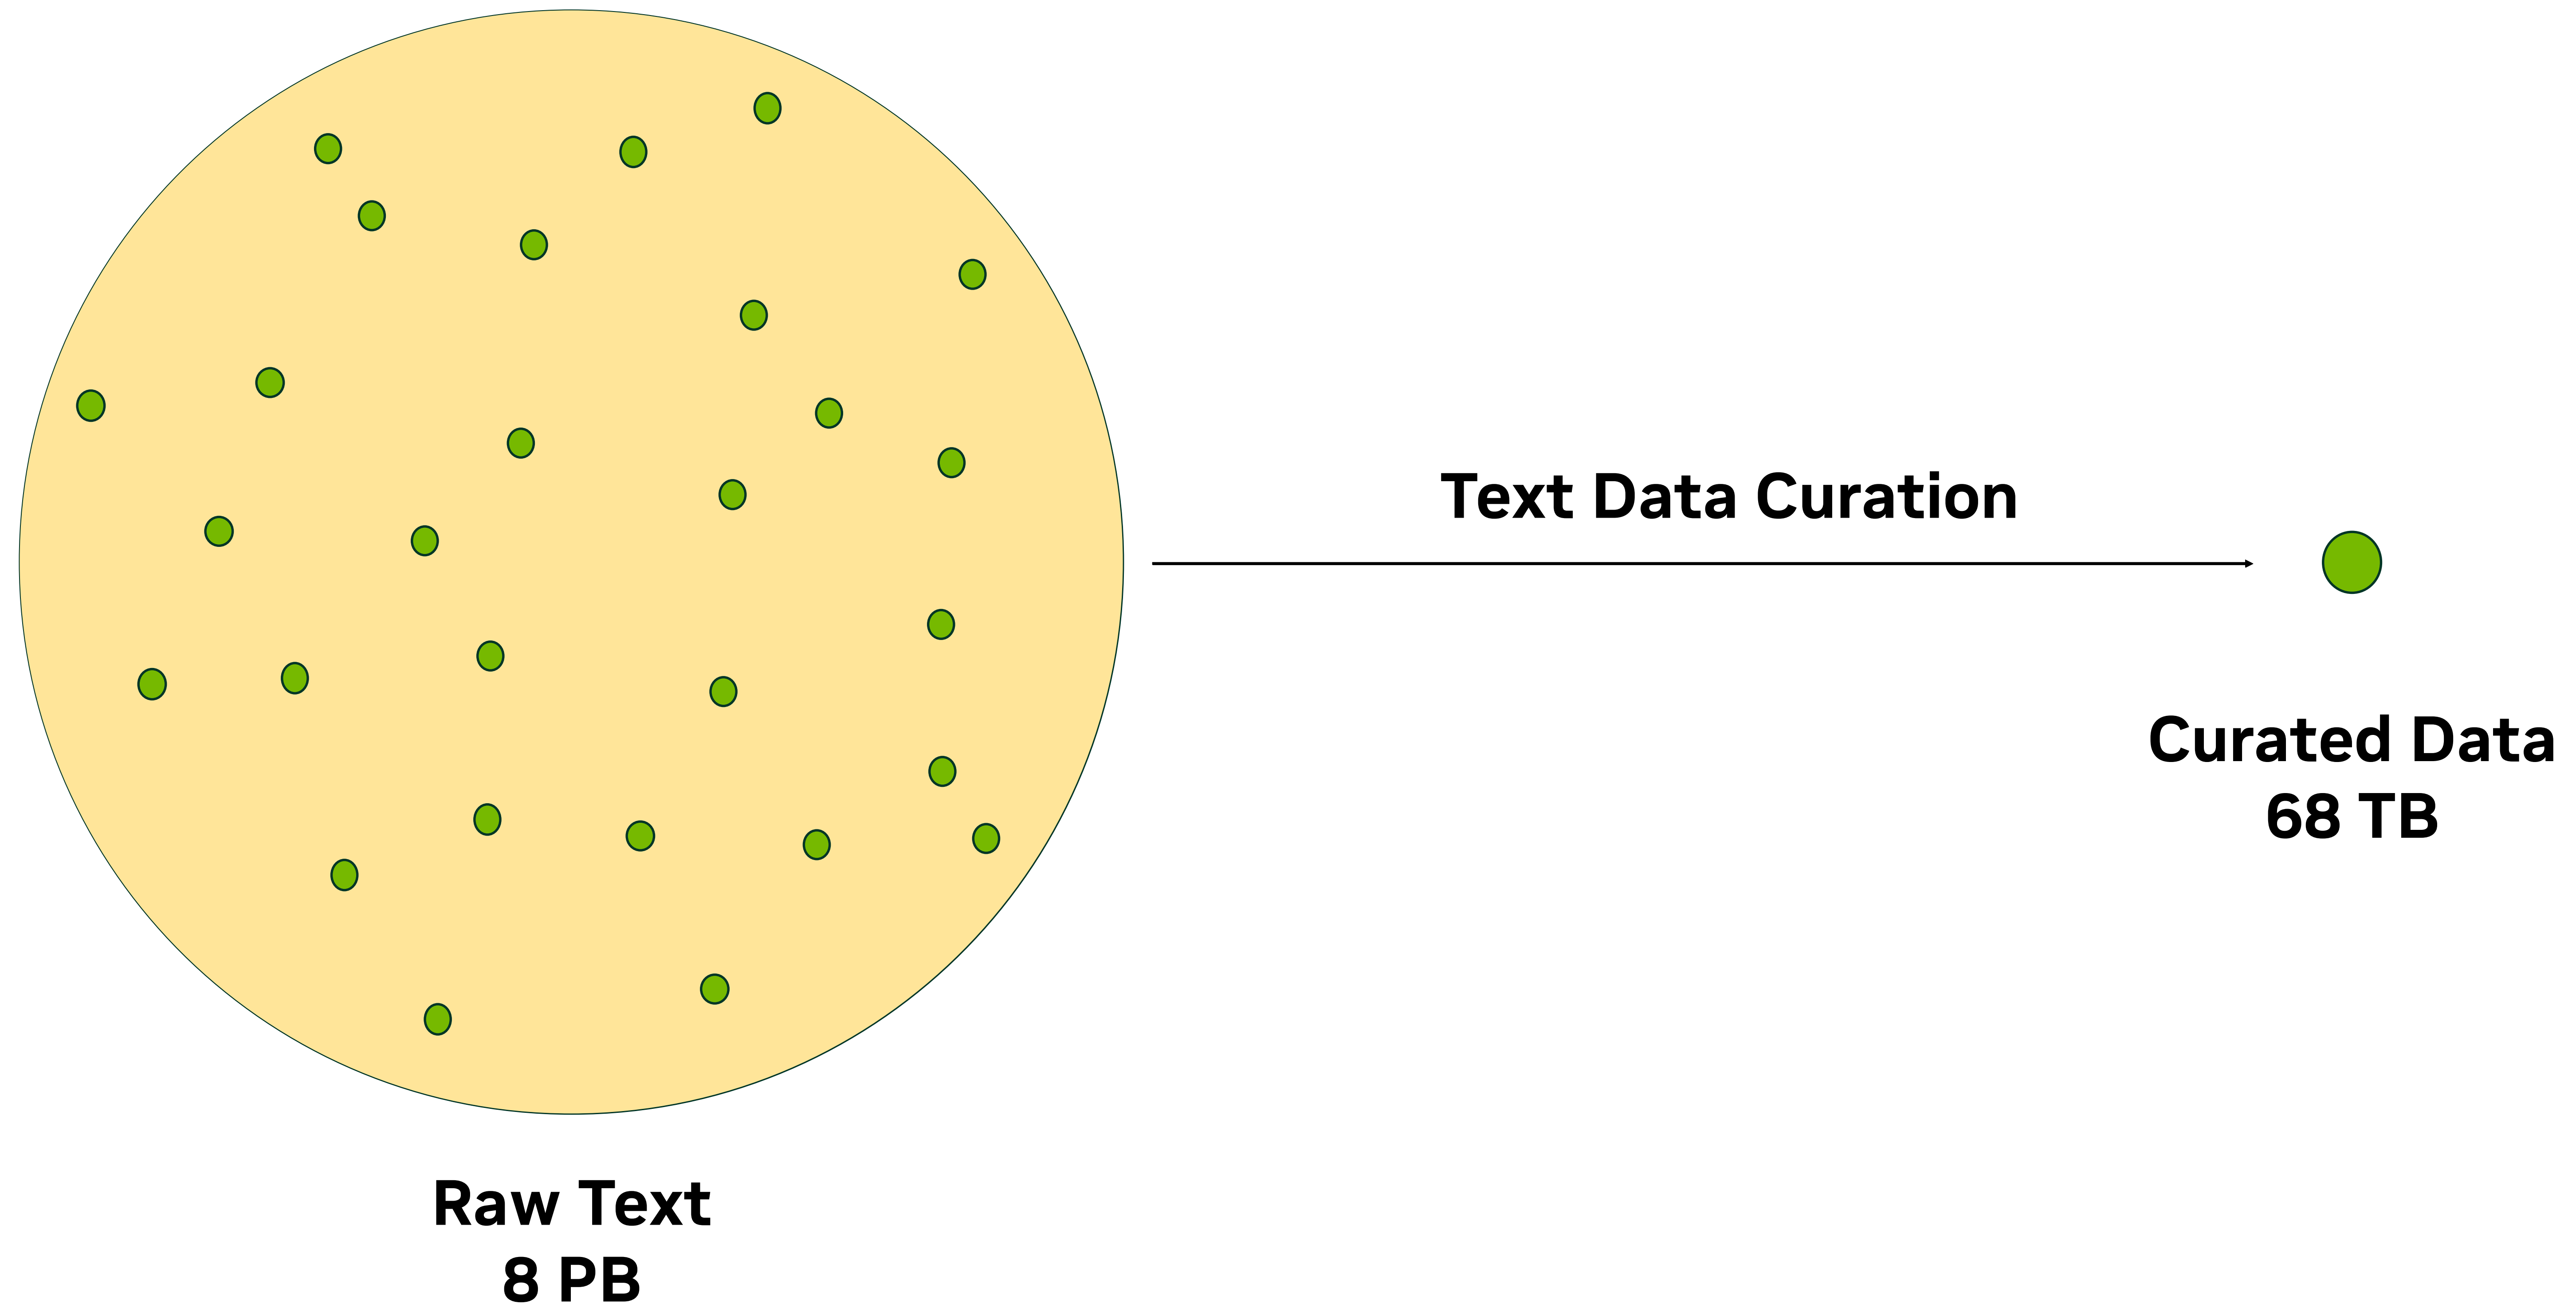
Reduces compute requirements by orders of magnitude, enabling faster iterations

Task Specialization




Enables optimization for specific tasks such as reasoning rather than general-purpose solutions

Only 1% of Raw Text Data is Curated to Train Foundation Models



Introducing: NeMo Curator for Text Processing

Easily integrate different features into your existing pipelines with Python APIs



[GitHub](#)

[NeMo framework container](#)

[PyPI](#)

NeMo Curator: Features to Train Foundation Models

Achieve higher accuracy with a variety of GPU-accelerated features



Synthetic Data Generation



Deduplication & Classification



GPU Acceleration with RAPIDS

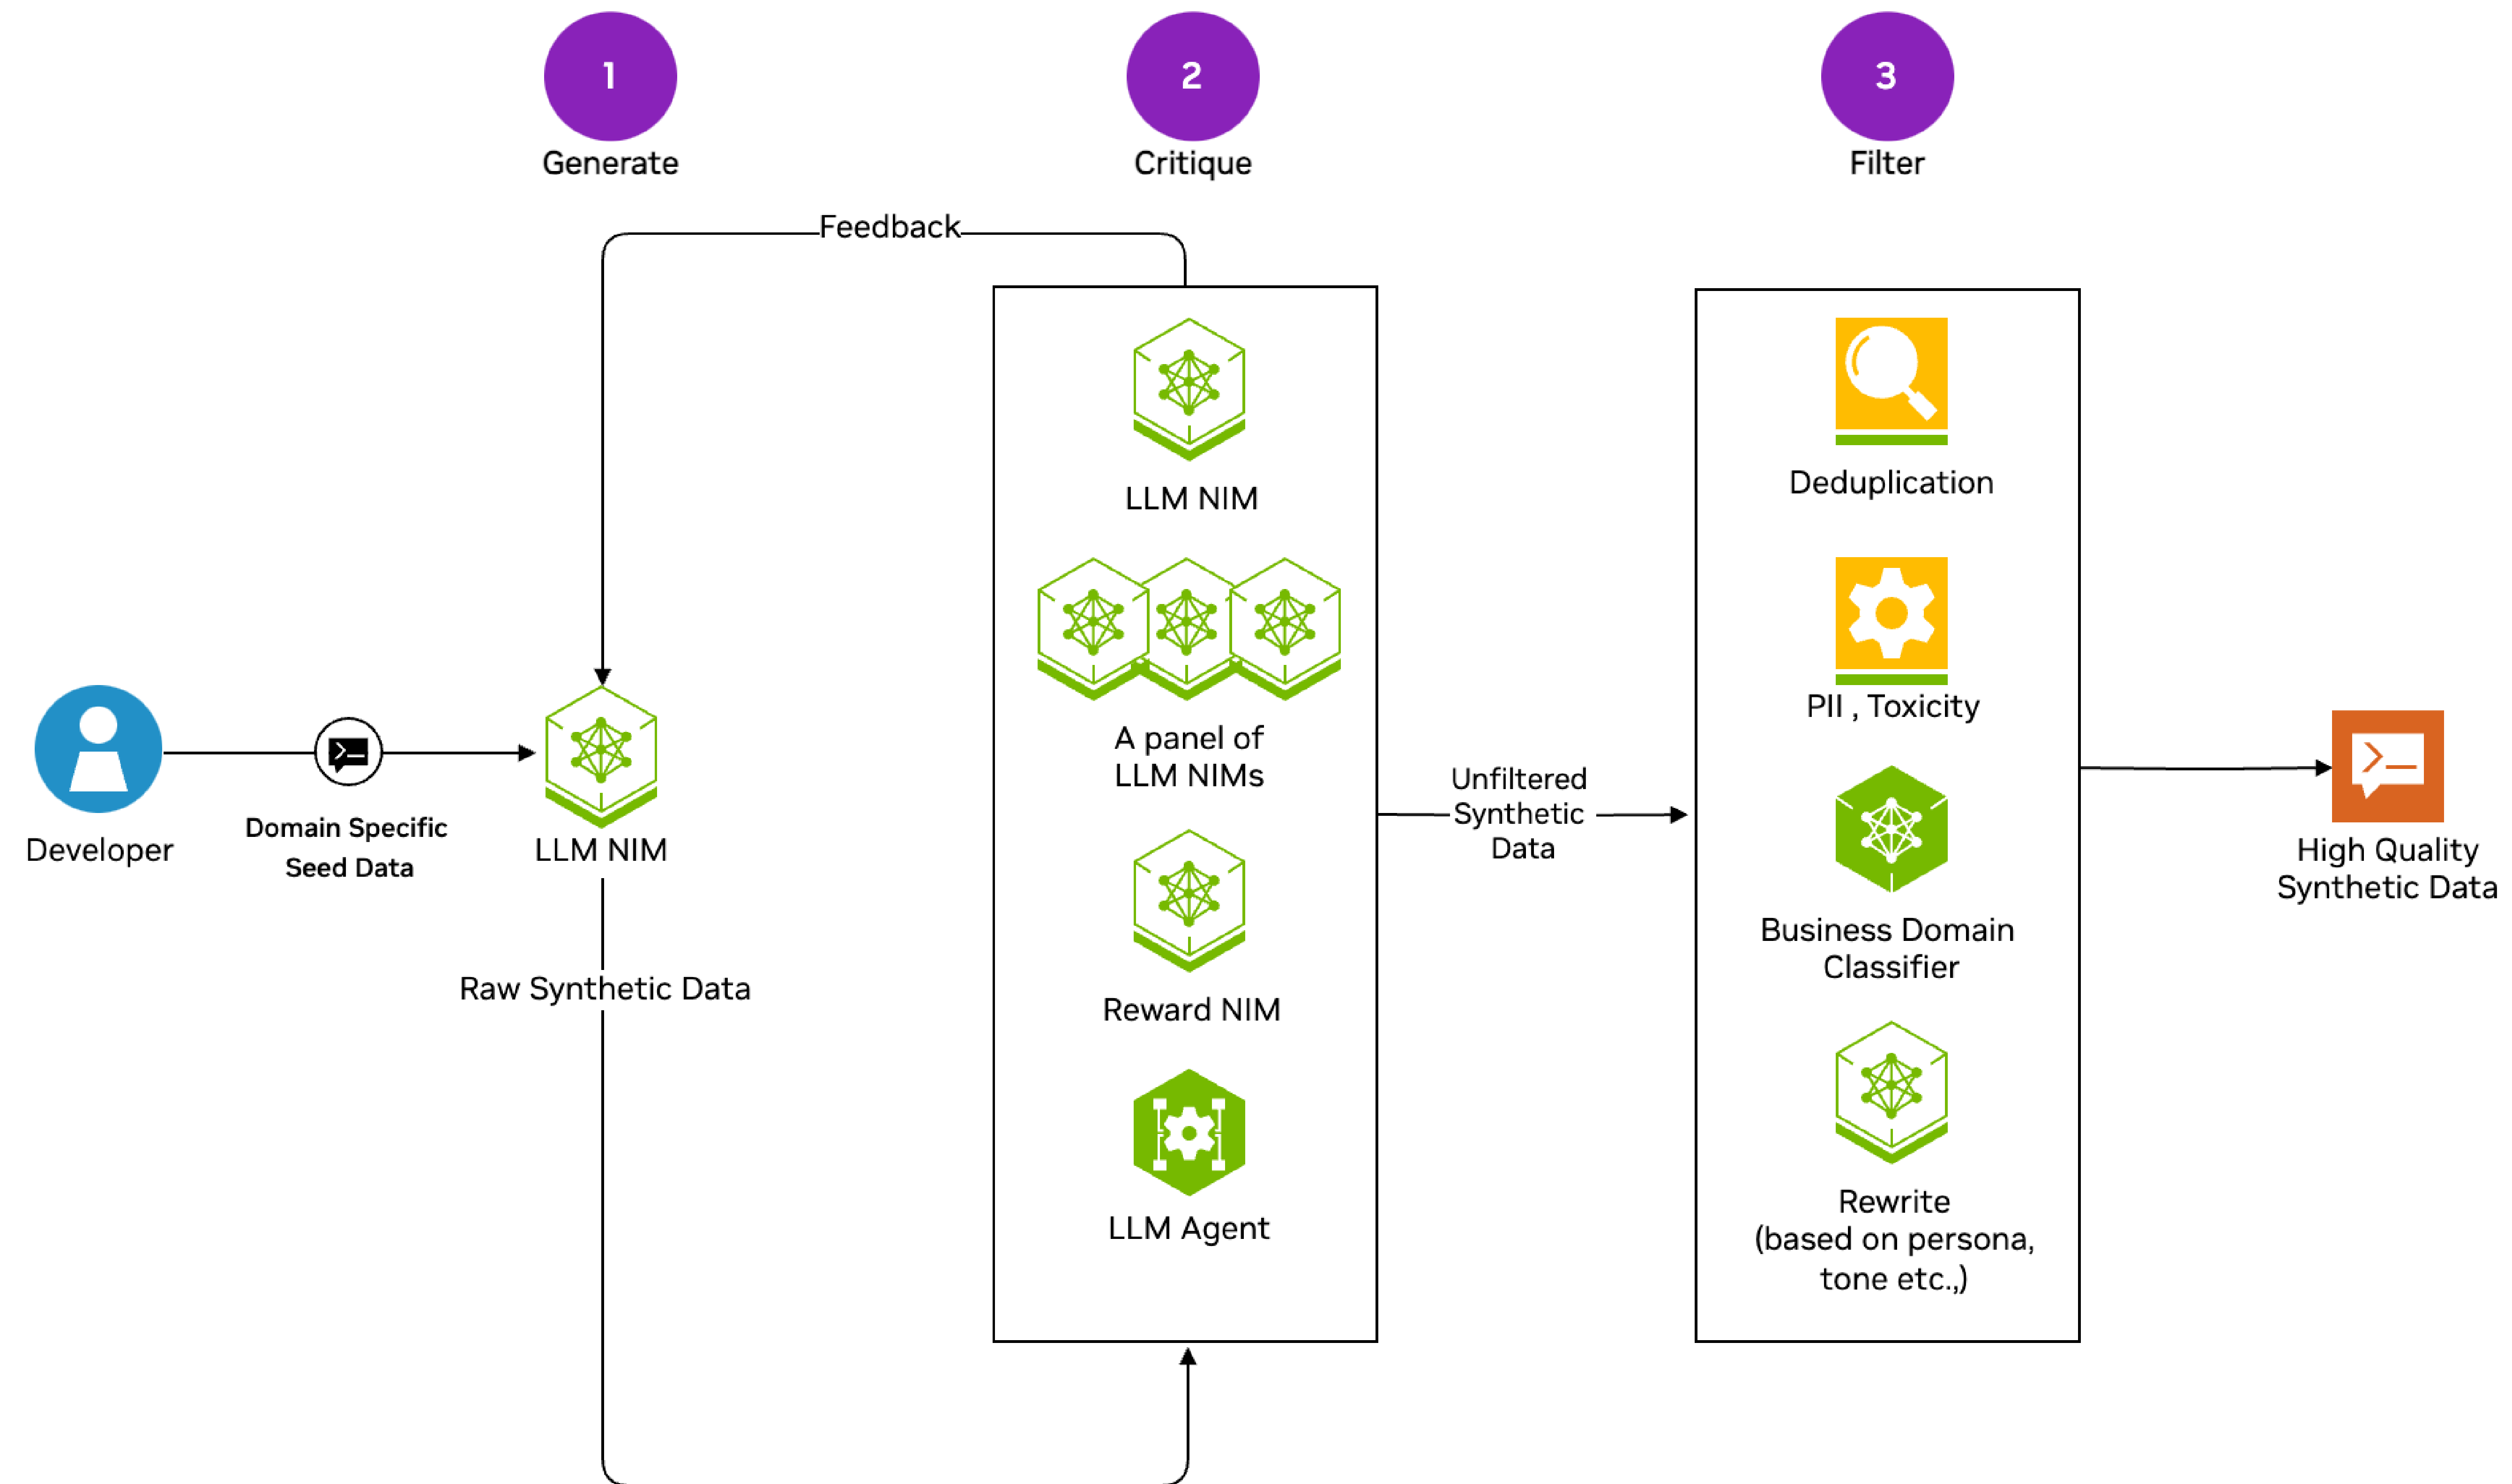
- **Pre-built pipelines** - for tasks like prompt generation, dialogue generation, and entity classification
- **Modular** - Easily integrate NeMo Curator's features into your existing pipelines
- **OpenAI API compatible** - Integrate custom Instruct and Reward models

- **Lexical Deduplication** – Identical (Exact) or near identical (Fuzzy)
- **Semantic Deduplication** – focuses on the meaning rather than the exact text
- **Classifier Models** - State-of-the-art open models to either enrich or filter your data.

- **cuDF** - for deduplication & classifier models
- **cuML** - for K-means clustering in semantic deduplication
- **cuGraph** – for fuzzy deduplication

Synthetic Data Generation

Easily get started with pre-built pipelines or integrate various features into your existing workflows



Synthetic Data Generation

Easily get started with pre-built pipelines or integrate various features into your existing workflows

Pre-built Pipelines

- Prompt generation (open q/a, closed q/a, writing, math/coding)
- Synthetic two-turn prompt generation
 - Dialogue generation
 - Entity classification

Llama 3.1 Nemotron 70B/340B Reward Model

- Benchmark-Topping Performance
- Single scalar scoring for human preferences
- Permissive license for commercial use

Tooling Support

- Integrate into existing pipelines
- Bring your custom Instruct/Reward Model
- Supports different filtering techniques

Synthetic Data Generation: Example

Use prompts to synthetically generate the QnAs and the reward score

```
model = "nvdev/nvidia/llama-3.1-nemotron-70b-instruct"

question_lst = []

for i in range(len(ragas)):
    closed_qa_responses = client.query_model(
        model=model,
        messages = [
            {
                "role": "user",
                "content": f"Generate a single, concise question similar to '{ragas.question[i]}' based on the provided context.\n"
                           f"Do NOT include any explanations, rationale, or additional text in your response—just the question itself:\n"
                           f"{ragas.contexts[i]}"
            }
        ],
        temperature = 0.2,
        top_p = 0.7,
        max_tokens = 1024,
    )

    question = closed_qa_responses[0]
    question_lst.append(question)
```

```
ragas['sdg_nemotron_q'] = question_lst
```

✓ 8.7s

```
ragas['sdg_nemotron_q']
```

✓ 0.0s

```
0  What is the significance of including a module...
1  How does the task execution stage in HuggingGP...
2  What are the primary technical hurdles encount...
3  How does Algorithm Distillation (AD) leverage ...
4  What are the core architectural elements and o...
5  How do consistent naming conventions, intra-fi...
6  How does API-Bank assess LLMs' decision-making...
7  How does a package.json fit into defining a No...
8  How does API-Bank assess LLMs' utilization of ...
9  How do finite context lengths and long-term pl...
Name: sdg_nemotron_q, dtype: object
```

```
reward_lst = []

model = "nvdev/nvidia/llama-3.1-nemotron-70b-reward"

for i in range(len(ragas)):
    messages = [
        {
            "role": "user",
            "content": f""
                       Provide a concise and well-reasoned answer to the following question:
                       "{ragas.sdg_nemotron_q[i]}"
                       Use the context below to ensure accuracy and clarity:
                       {ragas.contexts[i]}
                       ""
        },
        {
            "role": "assistant",
            "content": ragas.sdg_nemotron_a[i],
        },
    ]

    rewards = client.query_reward_model(messages=messages, model=model)
    rewards = float(rewards)

    reward_lst.append(rewards)
```

```
ragas['sdg_nemotron_reward'] = reward_lst
```

✓ 3.1s

```
ragas['sdg_nemotron_reward']
```

✓ 0.0s

```
0  -13.93750
1  -4.09375
2  -8.68750
3  -8.06250
4  -5.15625
5  -15.75000
6  -8.18750
7  -9.31250
8  -7.56250
9  -11.06250
Name: sdg_nemotron_reward, dtype: float64
```


SDG Demo Video

NeMo-Curator/tutorials/ymoslem/Law-StackExchlegal terms - What countNeMo/tutorials/llm/llam+

github.com/NVIDIA/NeMo-Curator/tree/main/tutorials/peft-curation-with-sdg

Presented by: **Mehran Maghoumi**

NeMo-Curator

Type to search

<> Code

Issues 38

Pull requests 13

Discussions

Actions

Projects 1

Security

Insights

main

NeMo-Curator / tutorials / peft-curation-with-sdg

Go to file

Add file

Maghoumi

[Tutorials] Synthetic data generation + reward model + curation for P...

9452189 · yesterday

History

Name	Last commit message	Last commit da...
..		
config	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
README.md	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
docbuilder.py	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
filters.py	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
main.py	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
modifiers.py	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday
synthetic_gen.py	[Tutorials] Synthetic data generation + reward model + curation for P...	yesterday

README.md

Curating Datasets for Parameter Efficient Fine-tuning with Synthetic Data Generation

This tutorial demonstrates the usage of NeMo Curator's Python API data curation as well as synthetic data generation, and qualitative score assignment to prepare a dataset for parameter-efficient fine-tuning (PEFT) of LLMs.

We demonstrate the pipeline using the [Law StackExchange dataset](#), which is a dataset of legal question/answers. Each record consists of a question, some context as well as human-provided answers.

In this tutorial, we implement various filtering and processing operations on the records. We then demonstrate the usage of external LLM services for synthetic data generation and reward models to assign qualitative metrics to each synthetic record. We further NeMo Curator's facilities to iteratively augment and refine the data until the dataset has reached the desired size.

Note: The use of external LLM services for synthetic data generation is entirely optional. Similarly, this tutorial can be executed on

Synthetic Data Generation: Example

Build an SDG pipeline to generate the data

main.py

```
14
15 > import argparse...
38
39 SCRIPT_DIR_PATH = os.path.dirname(os.path.abspath(__file__))
40 DATA_DIR = os.path.join(SCRIPT_DIR_PATH, "data")
41 TEMP_DIR = os.path.join(SCRIPT_DIR_PATH, "_temp")
42 CONFIG_DIR = os.path.join(SCRIPT_DIR_PATH, "config")
43 DATASET_URL = "https://huggingface.co/datasets/ymoslem/Law-StackExchange/resolve/main/law-stackexchange-questions-answers.json"
44
45
46 > def pre_imports():--
48
49
50 > def random_split_rows(rows: List[Any], train_ratio: float, val_ratio: float, seed=42):--
72
73
74 > def download_and_convert_to_jsonl() -> str:--
114
115
116 > def semantic_dedupe(dataset):--
144
145
146 > def run_curation_pipeline(--
237
238
239 > def run_pipeline(args, jsonl_fp):--
364
365
366 > def main():--
432
433
434 > if __name__ == "__main__":--
```

syn_gen.py

```
14
15 > import asyncio...
26
27 > PROMPT_GENERATE_QUESTIONS_FROM_ANSWER = """TEXT: --
36
37 > PROMPT_PARAPHRASE_TEXT = """TEXT: --
46
47
48 class SyntheticGenerator:
49
50 > def __init__(--
81
82 > def run(--
106
107 > def _split_sdg_responses(self, sdg_response: str) -> List[str]:--
120
121 > def _write_all_to_file(self, gen_entries, out_fp: str):--
182
183 > async def _prompt_model(--
252
253 > async def _synthesize_from_source(--
```

multiple rounds of SDG

```

  data
  curated
  final
  {} law-qa-test.jsonl
  {} law-qa-train.jsonl
  {} law-qa-val.jsonl
  round-1
  {} law-qa-train-synth-round-1.jsonl
  {} law-qa-train.jsonl
  round-2
  round-3
  round-4
  round-5
  raw
  downloads
  {} law-stackexchange-questions-answers.json
  splits
  {} law-qa-test.jsonl
  {} law-qa-train.jsonl
  {} law-qa-val.jsonl
```

one example on SDG curated result

```
{
  "id": "appliance-qa-B00074TB9U-synth-0",
  "asin": "B00074TB9U-synth-0",
  "question": "How effective is a ductless installation in eliminating cooking smells?",
  "answer": "The vent has two speeds; the lower one is suitable for regular air removal, while the higher speed is necessary for heavy-duty cooking. The higher speed functions well.",
  "questionType": "yes/no",
  "score": -1,
  "helpfulness": 0.6015625,
  "correctness": 0.388671875,
  "coherence": 3.078125,
  "complexity": 0.8671875,
  "verbosity": 0.60546875
}
```

Deduplication and Filtering

Scale across multi-node, multi-GPU setups, eliminating the need for iterative CPU processing

- Supports lexical and semantic deduplication for document processing
- Scales on multi node, multi-GPU by leveraging RAPIDS
- Scale to 100+ TB of data


Lexical Deduplication

- Utilizes text similarity to discover duplication
- **Exact deduplication**
 - Hashing based matching for each document
- **Fuzzy deduplication**
 - Minhash, Bucketization and Clustering

Semantic Deduplication

- Utilizes meaning of text to discover duplication
- Leverages embedding models to identify semantic documents
- Modularity to use custom embedding model

Acceleration



NeMo Curator - Deduplication

Scale across multi-node, multi-GPU setups, eliminating the need for iterative CPU processing

Value Proposition

- Accelerate on multi-node multi-GPU setups by leverage RAPIDS
- Scale up to 100+ TBs of data
- Get support for both
 - Lexical Deduplication: Uses text similarity, both Exact and Fuzzy
 - Semantic Deduplication: Uses text meaning

Technical Details

Lexical Deduplication

- Exact deduplication
 - Hashing for each document
- Fuzzy deduplication
 - Minhash, Bucketization and Clustering

Semantic Deduplication

- Leverage embeddings to identify semantic documents
- Use SOTA NeMo Retriever Text Embedding NIM or customize with your embedding model

NeMo Curator - Classifier Models

Create high-quality data blends with RAPIDS accelerated inference

Value Proposition

- Accelerated inference with RAPIDS powered distributed data classification module and intelligent batching
- Seamless scalability for classifying TBs of data
- Faster processing with parallelization across multiple GPUs
- Lower memory and compute footprint with state-of-the-art open classifier models (Apache 2.0 license)
- 8 classifier models released

Classifier Models

Domain Classifier

- Supports 26 domain classes
 - Top 10 classes : Finance, Health, Business and Industrial, Science, Law and Government, Internet and Telecom, Jobs and Education, News, Computers and Electronics, Shopping;
- Trained on 1 million Common Crawl samples & 500k Wikipedia articles
- Also available for multiple languages

Quality Classifier

- Classify quality of the document to 'High', 'Medium' or 'Low'
- Training data annotated by humans on quality factors such as: content accuracy, clarity, coherence, grammar, depth of information and overall usefulness of the document

NeMo Curator - Classifier Models

Create high-quality data blends with RAPIDS accelerated inference

Prompt Task and Complexity Classifier

- Classify tasks across 11 categories
 - OpenQA, Closed QA, Summarization, Text Generation, Code Generation, Chatbot, Classification, Rewrite, Brainstorming, Extraction, Other
- Compute complexity on 6 dimensions
 - Creativity, Domain Knowledge, Reasoning, Constraints, Contextual Knowledge, # of Few shots

FineWeb Nemotron-4 Edu Classifier

- Determine the educational value of a piece of text (score 0-5 from low to high).
- Trained using annotations from Nemotron-4-340B-Instruct

Content Type Classifier


- Categorize documents into 11 content types
 - Explanatory articles, News, Blogs, Boilerplate content, Analytical exposition, Online Comments, Reviews, Books & literature, Conversational, and Personal Websites.
- Useful for content management systems, digital publishers, recommendation systems

Instruction Data Guard

- Identify malicious prompts used for fine-tuning
- Optimal use for instruction:response datasets

Performance – Fuzzy Deduplication


~16x Faster Processing



Processing time for fuzzy deduplication of RedPajama-v2 subset (8TB/1.78T tokens)

'On': Data processed with NeMo Curator on 3 H100 nodes

'Off' : Data processed with a leading alternative library on CPUs



Processing time for fuzzy deduplication of RedPajama-v2 subset (8TB/1.78T tokens)

Scaling on 1, 2, 3, 4 H100 nodes 80GB

Deep Dive: Fuzzy Deduplication

Identifying similar documents

Document: "We the People of the United States, in Order to form a more perfect Union, establish Justice..."

Exact Duplicate: "We the People of the United States, in Order to form a more perfect Union, establish Justice..."


Fuzzy Duplicate: "Here is the US Constitution. We the People of the United States, in Order to form a more perfect Union, ensure Justice..."

Fuzzy Duplicate: "We the People of the United States, in Order to ensure Justice and domestic tranquility..."

Min-hashing

Doc -1					
K-shingle	{The quick brown}	{fox jumps over}	{the lazy dog}		
hashed-shingle	345L	3455L	934L		
Hash_1	23	49	50	23	Signature-1
Hash_2	56	24	39	24	
Hash_3	38	56	84	38	
Hash_4	48	29	93	29	
Hash_5	67	75	59	59	
Doc -2					
K-shingle		
hashed-shingle		
Hash_1	34	Signature-2
Hash_2	56	
Hash_3	78	
Hash_4	23	
Hash_5	14	
Doc -3					
...					


Signatures over hash buckets



Deep Dive: Fuzzy Deduplication

Min-Hashing and Bucketization


Document 1: “We the People of the United States, in Order to form a more perfect Union, establish Justice, ensure domestic Tranquility, provide for the common defense, promote the general Welfare, and secure the Blessings of Liberty...”




Deep Dive: Fuzzy Deduplication

Min-Hashing and Bucketization


Document 2: “Here is the US Constitution. We the People of the United States, in Order to form a more perfect Union, establish Justice, insure foreign Tranquility, provide for the common defense, promote the general Welfare, and secure the Blessings of Liberty...”




Bucket #1



Bucket #2



Bucket #3




Bucket #4


Deep Dive: Fuzzy Deduplication

Min-Hashing and Bucketization


Document 3: “We the People of Canada, in Order to form a more perfect Union, establish Justice, insure domestic Tranquility, provide for the common defense, promote the general Welfare, and secure the Blessings of Liberty...”




Bucket #1



Bucket #2



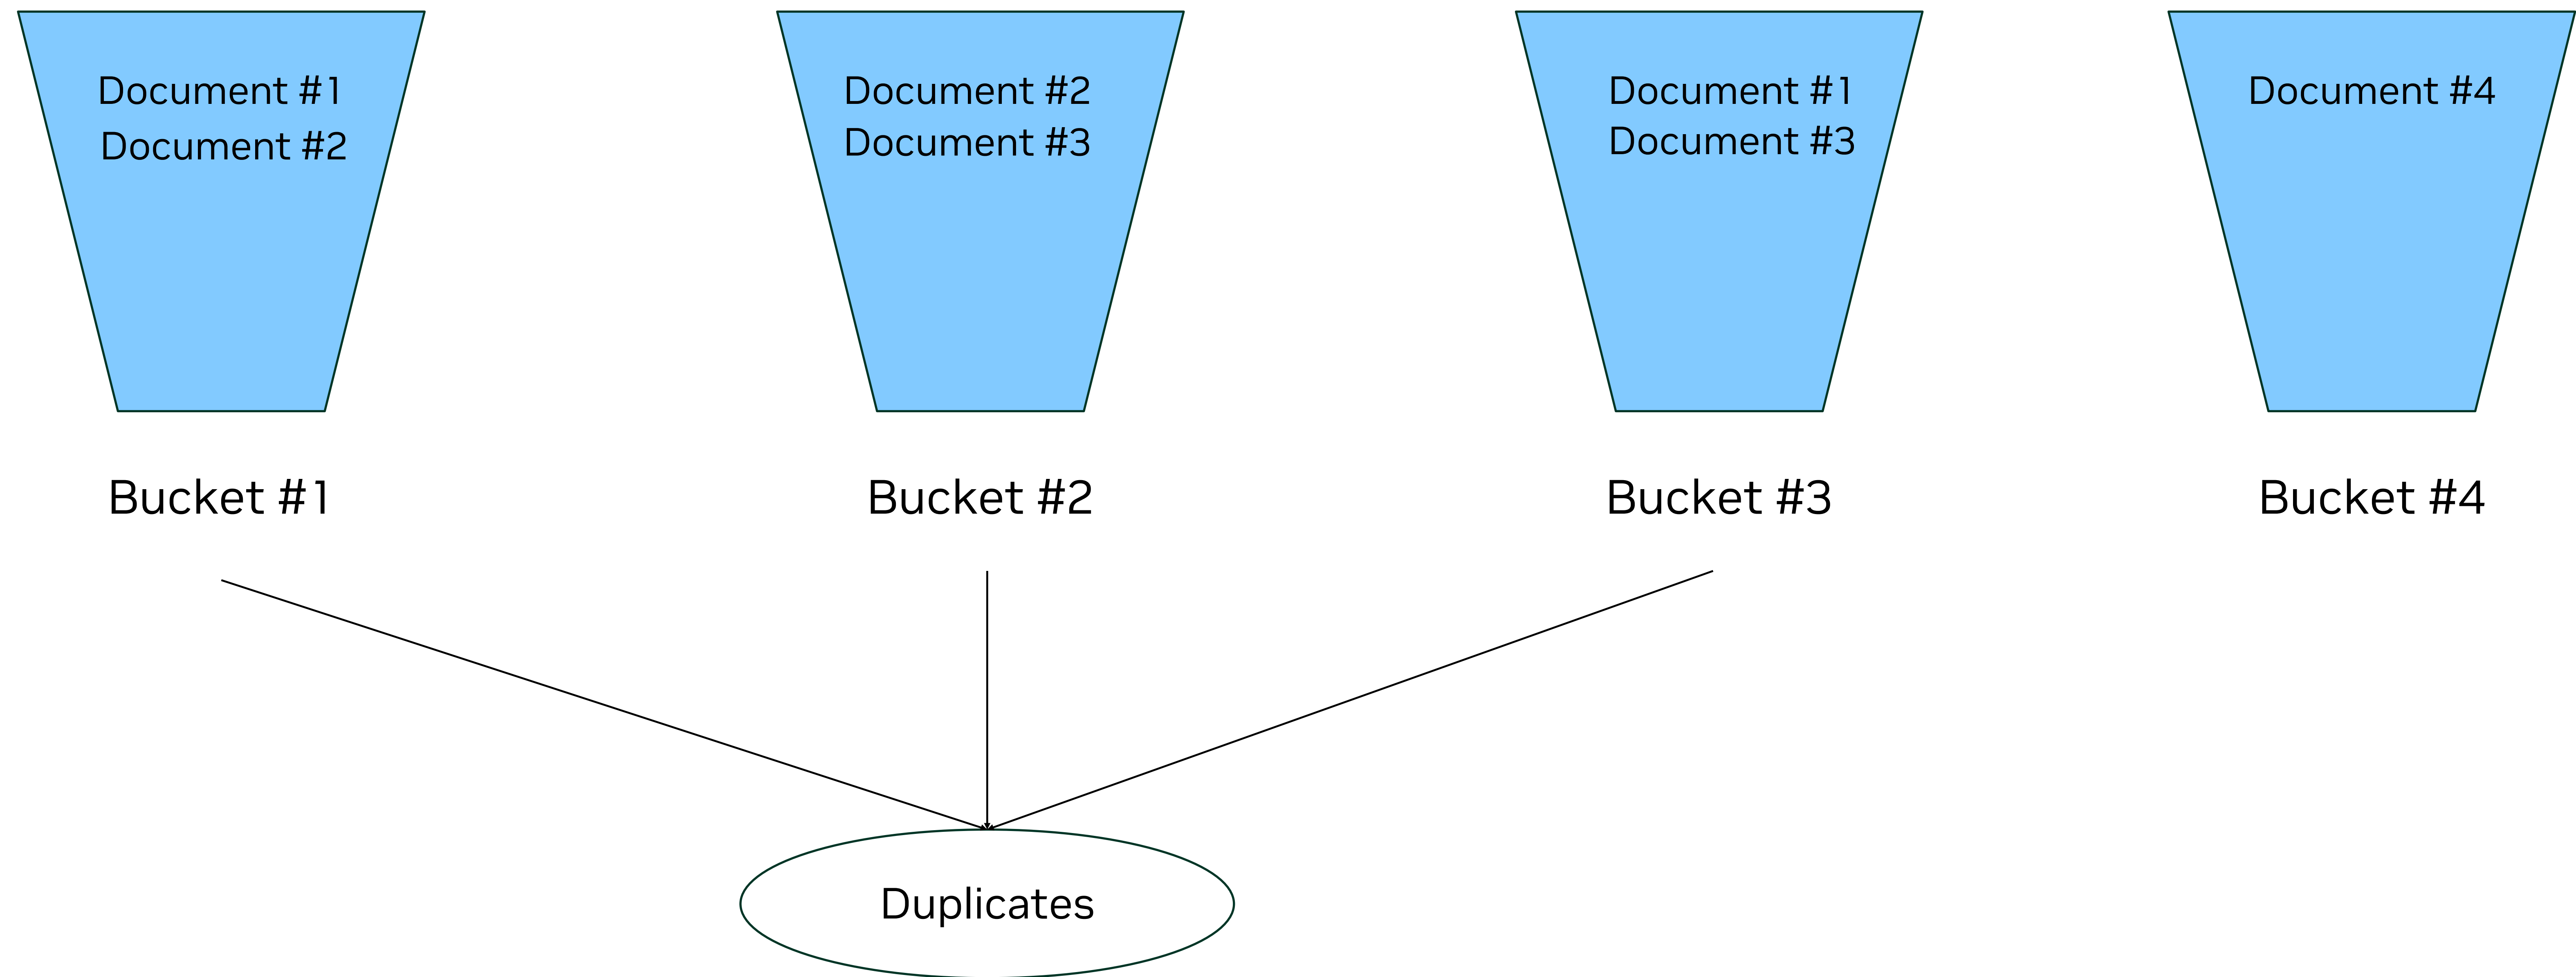
Bucket #3



Bucket #4

Deep Dive: Fuzzy Deduplication

Min-Hashing and Bucketization



Semantic Deduplication: Example

Remove redundant data by identifying and eliminating semantically similar data points

sem_dedup_config.yaml

```
# Configuration file for semantic dedup
cache_dir: "semdedup_cache"
num_files: -1

# Embeddings configuration
embeddings_save_loc: "embeddings"
embedding_model_name_or_path: "sentence-transformers/all-MiniLM-L6-v2"
embedding_batch_size: 128

# Clustering configuration
clustering_save_loc: "clustering_results"
n_clusters: 1000
seed: 1234
max_iter: 100
kmeans_with_cos_dist: false

# Semdedup configuration
which_to_keep: "hard"
largest_cluster_size_to_process: 100000
sim_metric: "cosine"

# Extract dedup configuration
eps_thresholds:
  - 0.01
  - 0.001

# Which threshold to use for extracting deduped data
eps_to_extract: 0.01
```


Classification and PII

Create high-quality data blends with RAPIDS accelerated inference

- Accelerated inference through distributed classification model and intelligent batching
- Redact/remove Personally Identifiable information (PII) using SOTA model


Classification

- **Domain Classifier**
 - Model trained on 1.5 million samples
 - Classifies text in 26 domain classes
- **Document Quality Classifier**
 - Classifier quality of document in 'High', 'Medium', 'Low' categories
 - Enables document quality check

Mask/remove PII


- Remove sensitive information from data
- Utilizes State-Of-The-Art spacy model to remove PII information

Acceleration




Accelerated Data Processing Maximizes LLM Performance & Scale

Nemo Curator: Up to 7% better Accuracy for LLM Downstream Tasks



Scale to 100+ TB of data



NeMo Curator: Example

```
# Download your dataset
dataset = download_common_crawl("/datasets/common_crawl/", "2021-04", "2021-10", url_limit=10)
# Build your pipeline
curation_pipeline = Sequential([
    # Fix unicode
    Modify(UnicodeReformatter()),
    # Discard short records
    ScoreFilter(WordCountFilter(min_words=80)),
    # Discard low-quality records
    ScoreFilter(FastTextQualityFilter(model_path="model.bin")),
    # Discard records from the evaluation metrics to prevent test set leakage.
    TaskDecontamination([Winogrande(), Squad(), TriviaQA()])
])
# Execute the pipeline on your dataset
curated_dataset = curation_pipeline(dataset)
```

[GitHub](#)[NeMo framework container](#)[PyPI](#)

NeMo Curator - Resources

Getting Started

- [NeMo Framework Container](#)
- [GitHub](#)
- [PyPI](#)
- [Installation Guide](#)
- [Developer Page](#)
- [User Guide/ Docs](#)
- [API Docs](#)
- [Classifier Models](#)
- [Examples](#)
- [Best Practices](#)
- [Bugs](#)
- [Discussions](#)

Tutorials & Blogs

Pre-training / DAPT

- [Curating data for LLM training \(Blog\)](#)
- [Curating non-English data \(Blog\)](#)
- [How-to run classifier models](#)
- [All blogs](#)

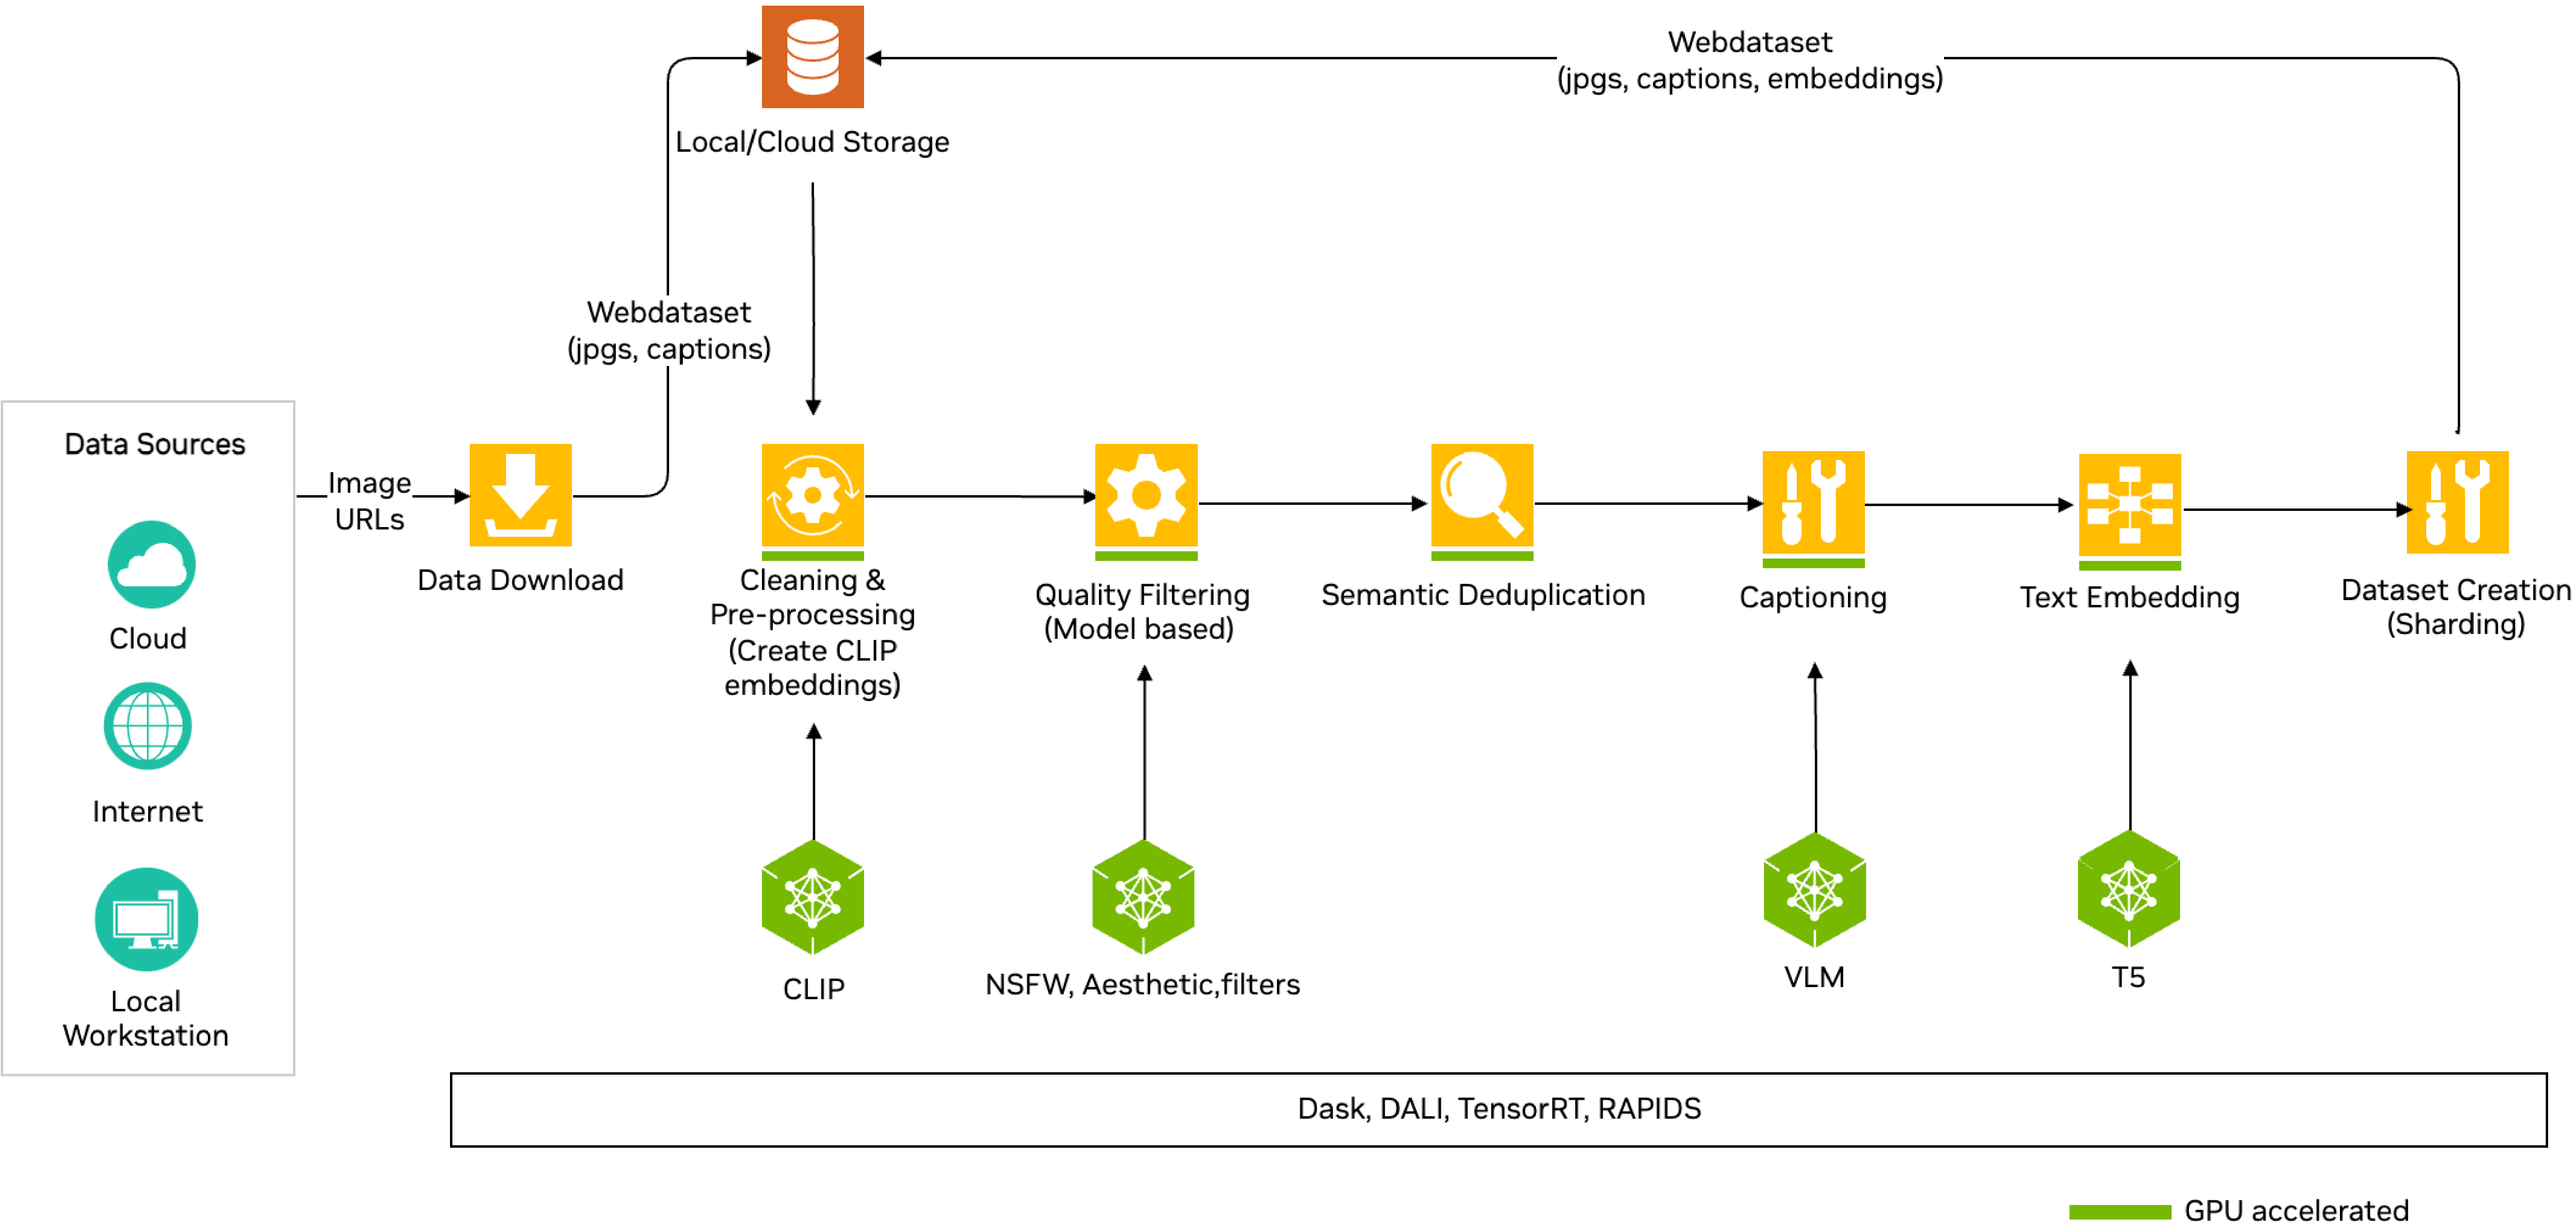
Fine-tuning

- [Curating data for PEFT \(Blog\)](#)
- [Curating synthetic data for PEFT](#)
- [SDG using Llama 3.1 405B & Nemotron 4-340B](#)
- [SDG using Nemotron 4-340B](#)



Image Processing

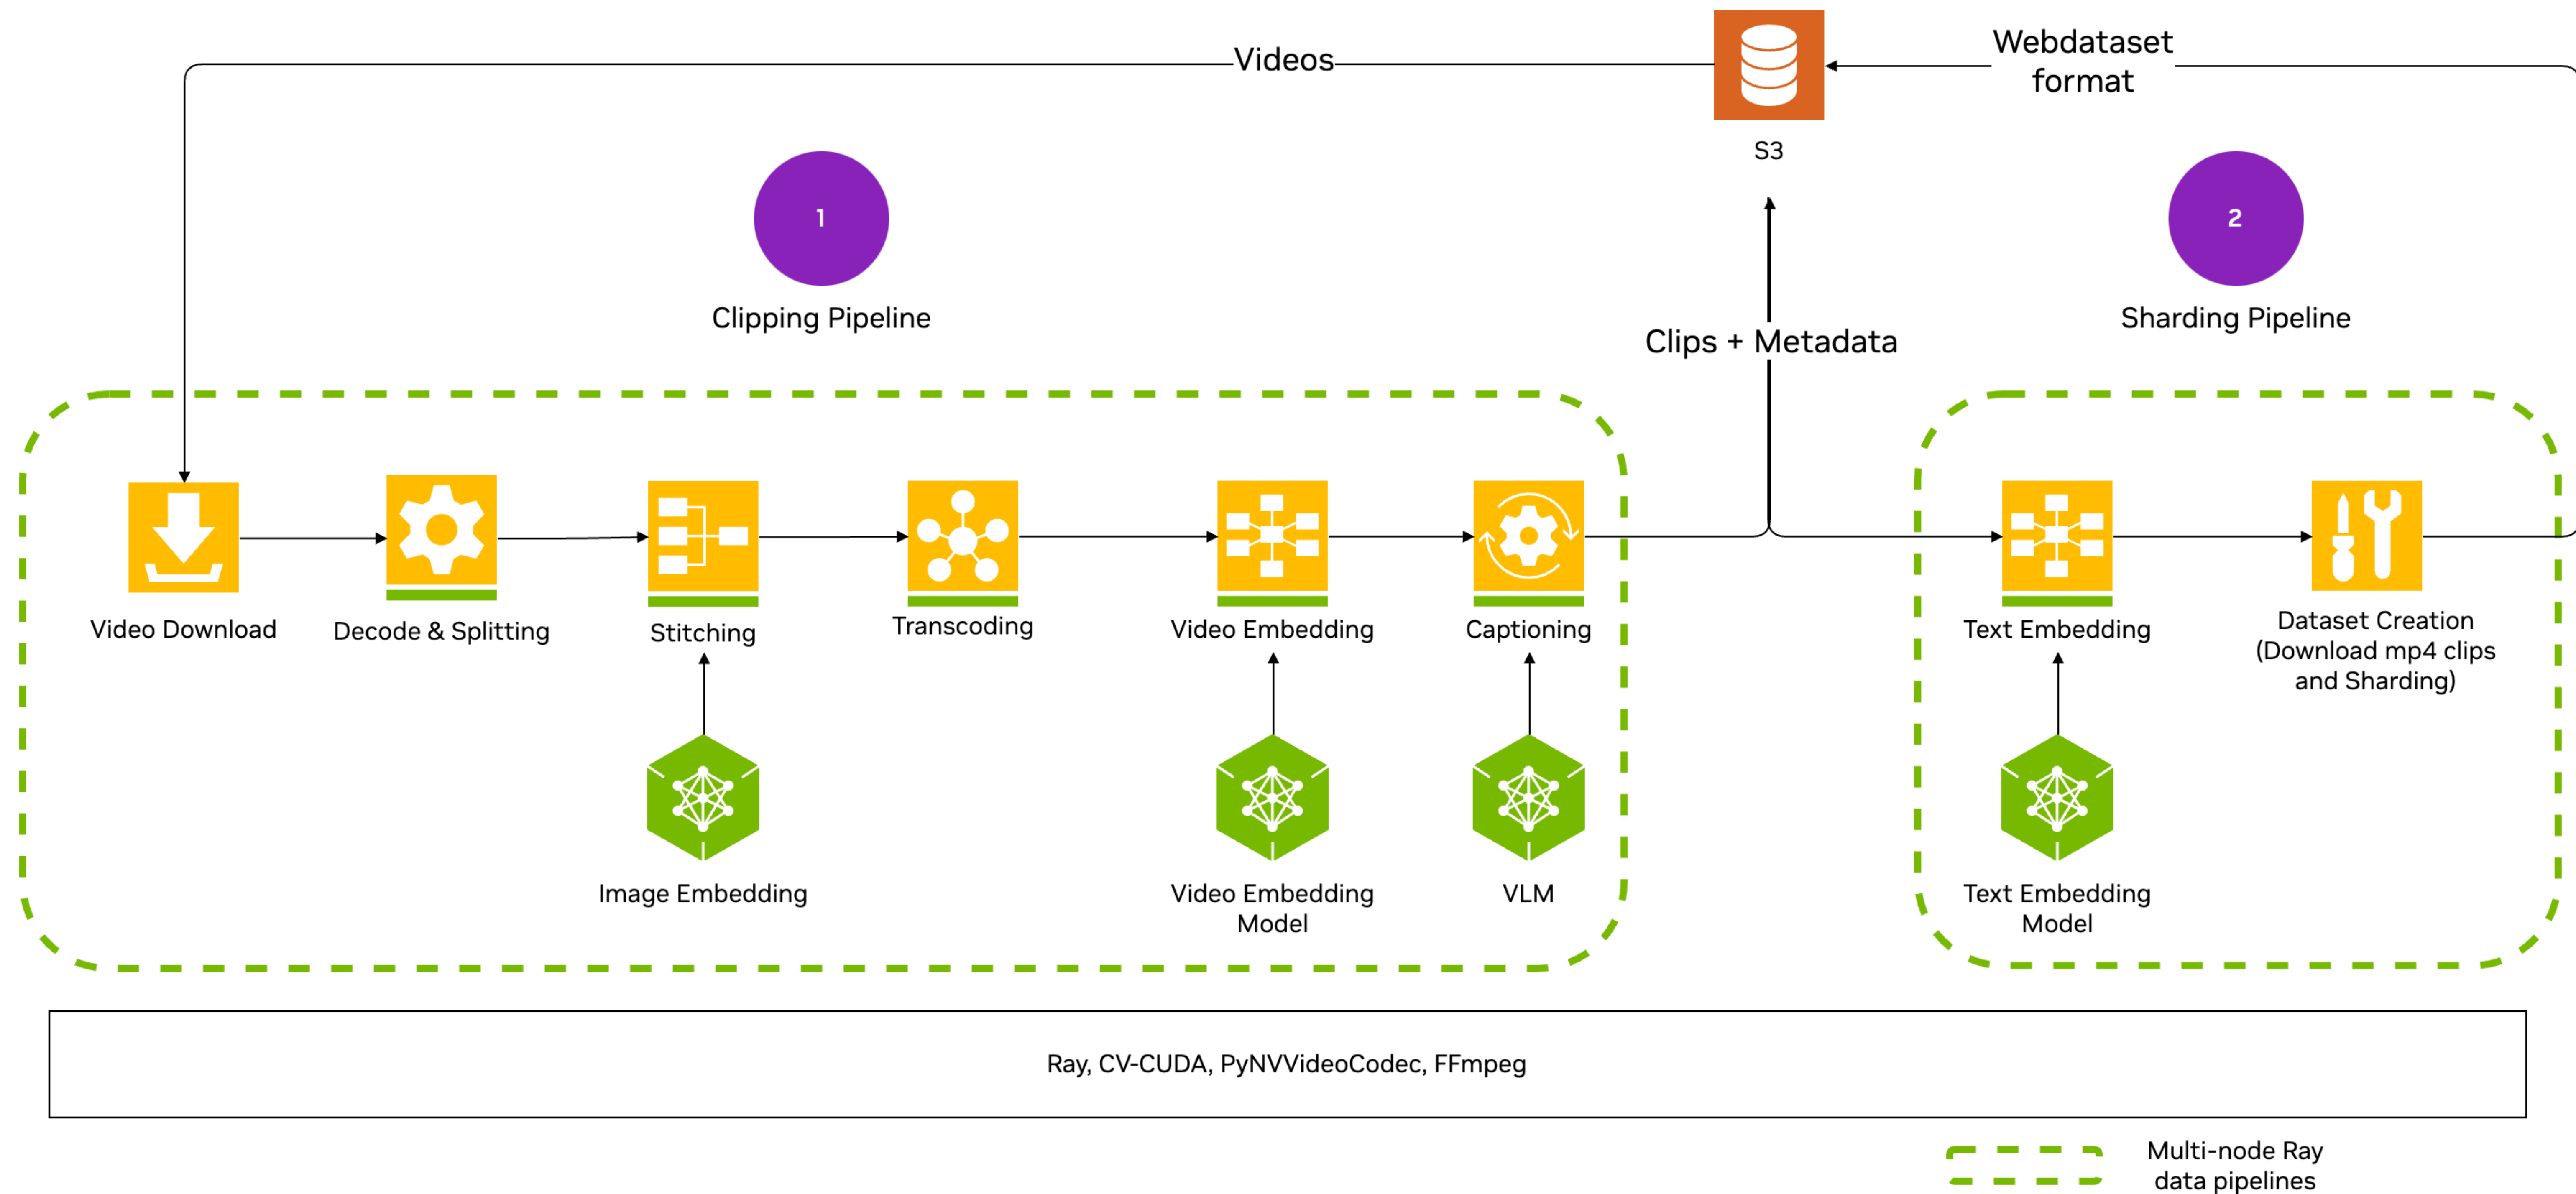
NeMo Curator Architecture: Image Processing





Video Processing

Video Curation: EA Features



State of the Art Models

Example Use case: Generate new lighting, weather, and geolocations




Best Practices

- Choosing the Right Quality Model Type
- Handling GPU Out-of-Memory (OOM) Errors
 - Controlling Partition Sizes
- Fuzzy Deduplication Guidelines
 - Reduce bucket counts
 - Reduce buckets per shuffle
 - Adjust files per partition
- GPU Memory and Utilization – Dask Dashboard

Developer Tools and Resources

Accelerate innovation and growth



Learn more: developer.nvidia.com

Individuals

Software

100s of APIs, models, SDKs, microservices, and early access to NVIDIA tech

Training

Hands-on self-paced courses, instructor-led workshops, and certifications

GPU Sandbox

Approval basis, multi-GPU and multi-node

Learning

Tutorials, self-paced courses, blogs, documentation, code samples

Community

Dedicated developer forums, meetups, hackathons

Ecosystem

GTC, NVIDIA Partner Network

Organizations

Startups

Cloud credits, engineering resources, technology discounts, exposure to VCs

Venture Capital

Deal flow and portfolio support for Venture Capital firms

Higher Education

Teaching kits, training, curriculum co-development, grants

ISVs and SIs

Engineering guidance, discounts, marketing opportunities

Research

Grant programs, collaboration opportunities

Enterprises

Tailored developer training, skills certification, technical support



Thank You